

Towards Improved Provisioning and Utilization of Resources in Virtualized Environments

Sujesha Sudevalayam

Department of Computer Science and Engineering
Indian Institute of Technology Bombay
{sujesha}@cse.iitb.ac.in

January 12th, 2018
PhD Defence Presentation

Computing-as-a-Service: The New Norm



Electricity Grid



Public Transport



Virtualization
Containers
Edge Computing
Serverless Computing

Enabling technology

- Software as a Service
- Platform as a Service
- Infrastructure as a Service

Computing-as-a-Service: The New Norm



Electricity Grid



Public Transport



Enabling technology

Virtualization
Containers
Edge Computing
Serverless Computing

- Software as a Service
- Platform as a Service
- Infrastructure as a Service

1. Network-affinity aware CPU Usage Estimation

Prediction of virtualized CPU usage for inter-PM and intra-PM network communication between VMs

- *Affinity-aware Modeling of CPU Usage for Provisioning Virtualized Applications.* Proceedings of the 4th International Conference on Cloud Computing (CLOUD), 2011. Sujesha Sudevalayam and Purushottam Kulkarni.
- *Affinity-aware Modeling of CPU Usage with Communicating Virtual Machines.* Journal of Systems and Software (JSS), 2013. Sujesha Sudevalayam, Purushottam Kulkarni.

Thesis Contributions

1. Network-affinity aware CPU Usage Estimation

2. VM Disk I/O Reduction by Host-cache Manipulation

Reduction of disk I/O by exploiting content similarity within and across virtual machines

- *DRIVE: Using Implicit Caching Hints to achieve Disk I/O Reduction in Virtualized Environments*. Proceedings of the 21st International Conference on High Performance Computing (HiPC), 2014. Sujesha Sudevalayam, Purushottam Kulkarni, Rahul Balani and Akshat Verma.

Thesis Contributions

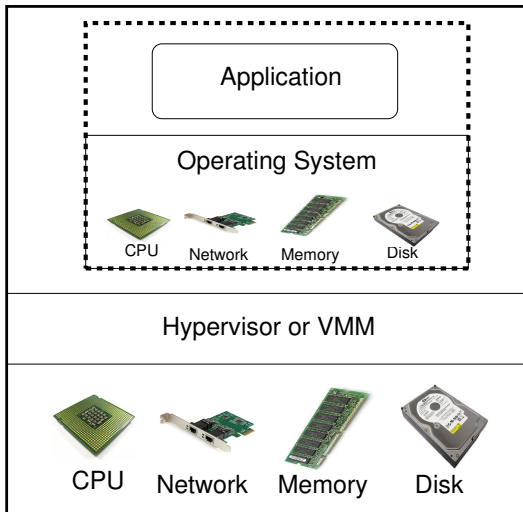
1. Network-affinity aware CPU Usage Estimation

2. VM Disk I/O Reduction by Host-cache Manipulation

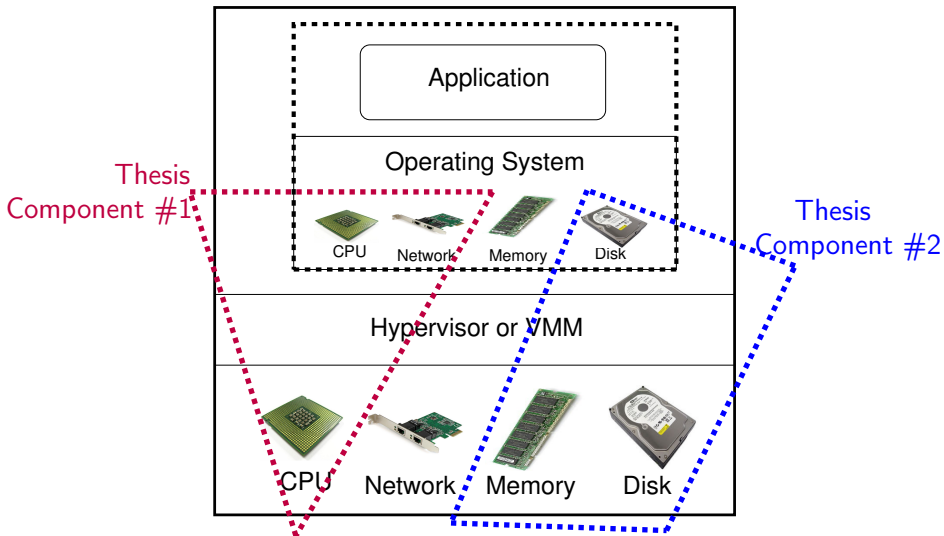
Tools and Deliverables

- ① WLoadGen: A load generator for CPU, disk & network loads
- ② SimReplay: A simulator for analyzing host cache effectiveness
- ③ preadwritedump: A kernel module for I/O request tracing

Resources Under Consideration In Virtualized Environment



Resources Under Consideration In Virtualized Environment



Thesis Contributions

1. Network-affinity aware CPU Usage Estimation

Prediction of virtualized CPU usage for inter-PM and intra-PM network communication between VMs

2. VM Disk I/O Reduction by Host-cache Manipulation


Reduction of disk I/O by exploiting content similarity within and across virtual machines

Network-affinity aware CPU Usage Estimation


- 1 Event profiling study showing difference between *intra-PM* and *inter-PM* network code paths
- 2 Benchmarking of CPU usage in colocated and dispersed scenarios for a VM pair
- 3 Pair-wise linear regression model to predict total CPU when network traffic changes nature between intra-PM and inter-PM
- 4 Pair-wise linear regression model to predict differential CPU usage
- 5 Application of pair-wise models to predict for multi-VM scenarios

Migration-Enabled Resource/Performance Management

Colocate VMs for Resource Efficiency => intra-PM network traffic




Disperse VMs for QoS => inter-PM network traffic



- Both collocation and dispersion need **resource usage estimation**
- Incorrect estimation is sub-optimal
 - Under-estimation => degraded performance
 - Over-estimation => wasted resources

Mutable and Immutable Network traffic for *Migratory* VMs




Definition of Mutable n/w traffic

For a VM pair, network traffic whose nature may *change between inter-PM and intra-PM*

Our hypothesis

Mutable network traffic has *different CPU overheads* in colocated and dispersed scenarios => ignoring affinity effects could result in incorrect CPU usage estimation

Communicating VMs (Xen-view)




- **Dom0 overhead** for DomU's I/O activity (network & disk)
- Intra-PM network traffic
 - Dom0 does not use native I/O drivers
 - Shared memory based copying of packets
- **Less CPU overhead for intra-PM** traffic compared to *inter-PM*
- Needs to be accounted for during VM migration


Effect of colocation on CPU usage for *Mutable* N/w traffic

Benchmarking setup: 2 VMs on 2 PMs—dispersed and colocated scenarios

Network load: Transmitted (Tx) by one VM and Received (Rx) by other



(a) Receiving DomU CPU util



(b) Dom0 CPU util for Rx/Tx

Observations

- **DomU:** Rx increase from 20-90 Mbps => decrease of 2-8% CPU util
- **Dom0:** Increase from 20-90 Mbps => decrease of 9-25% CPU util
- Linear correlation between network and CPU usage

Effect of colocation on CPU usage for Immutable n/w traffic, CPU and disk loads

Benchmarking setup: 4 VMs on 4 PMs—dispersed and colocated scenarios

Table : Percentage CPU usage for Immutable Rx

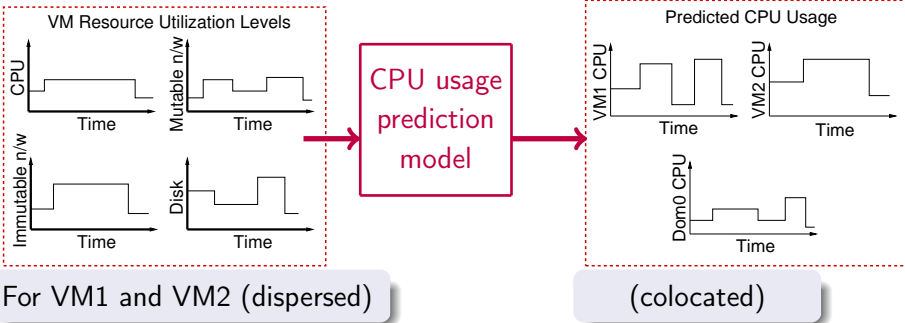
Immutable Rx (Mbps) ($< VM_1, VM_2 >$)	% CPU utilization	
	Dispersed case $VM_1, VM_2, \sum Dom0_i$	Colocated case $VM_1, VM_2, Dom0$
$<20, 50>$	4, 7, 18	4, 7, 14
$<40, 10>$	6 , 2 , 15	6 , 2 , 11
$<60, 10>$	8, 2, 18	8, 2, 14

Observations

- 1 No change in DomU CPU usage between colocated and dispersed
- 2 Dom0 CPU usage change of **4%** for extra Dom0 instance (constant)
- 3 Similar observations for other workloads—CPU and disk read/write

Problem: Affinity-aware Resource Requirement Estimation

Given a pair of VMs and their resource utilization levels, predict the CPU resource requirement of DomU & Dom0, when VM placement scenario changes between dispersed and colocated.



Core Idea

Build **linear prediction models** since correlation of CPU usage with all other resources usage is linear

Linear Regression Modeling for CPU Estimation

Parameters in the models

- **CPU** metrics: user, system, iowait
- **Disk** metrics: read blocks/second, write blocks/second
- **Mutable and immutable** network metrics: Rx and Tx Kbps, packets/sec

DomU Models

$$CPU_{colocated} = f(CPU, Disk, Mutable, Immutable)_{dispersed}$$


$$CPU_{dispersed} = f(CPU, Disk, Mutable, Immutable)_{colocated}$$

Dom0 Models


$$CPU_{colocated} = f(CPU_1, Disk_1, Mutable_1, Immutable_1, \\ CPU_2, Disk_2, Mutable_2, Immutable_2)_{dispersed}$$

$$CPU_{dispersed} = f(CPU_1, Disk_1, Mutable_1, Immutable_1)_{colocated}$$

Prediction for Synthetic workloads - Xen Dom0 model



(a) Dispersed to colocated




(b) Colocated to dispersed


Observations

90th percentile prediction error within 3% absolute CPU utilization, and maximum error 5-6% absolute CPU (Similarly for RUBiS workload as well)

Prediction for Synthetic workloads - Xen Dom0 model



(a) Dispersed to colocated





(b) Colocated to dispersed

Observations

90th percentile prediction error within **3% absolute CPU** utilization, and maximum error 5-6% absolute CPU (Similarly for RUBiS workload as well)

Building an Enhanced Prediction Model

Because “differential” CPU usage is only due to mutable n/w traffic



Evaluation of Differential CPU Prediction Models




Figure : Colocated DomU-1 (Synthetic)




Figure : Colocated DomU-1 (RUBiS)




Figure : Colocated Dom0 (Synthetic)




Figure : Colocated Dom0 (RUBiS)

Result

Maximum prediction error between 1-2% absolute CPU utilization.

Applying Pair-wise Models to Multi-VM Scenarios

PM1	Dom0	DomU1	DomU2
PM2	Dom0		
PM3	Dom0	DomU3	

"Combined" transition

PM1	Dom0	DomU1	
PM2	Dom0		
PM3	Dom0	DomU3	DomU2

Two-step prediction for combined transition

- 1 Predict using *dispersion model*
- 2 Predict using *colocation model* on previous prediction

Evaluated for multi-hop transitions

Table : Maximum error in Dom0 CPU utilization prediction

Transition	Max error (% absolute CPU)		
	Dom0-PM1	Dom0-PM2	Dom0-PM3
Transition (i)	0.75	-	-
Transition (ii)	1.99	-	0.85
Transition (iii)	-	0.51	0.43

Summary and Conclusions

- **Colocation of mutually-communicating VMs impacts their CPU requirement**
 - **DomU:** For Rx, increase from 20 to 90 Mbps => decrease from 2% to 8% CPU requirement
 - **Dom0:** Increase from 20 to 90 Mbps => decrease from 9% to 25% CPU requirement
- **Simple linear model shown to predict “differential” CPU requirement from mutable n/w traffic profiles**
 - **Synthetic and RUBiS workloads:** Max error within 1.5% absolute CPU utilization for both DomU and Dom0 models
 - **Multi-VM scenario:** Max error within 2% for all transitions

Publications

- *Affinity-aware Modeling of CPU Usage for Provisioning Virtualized Applications.* Proceedings of the 4th International Conference on Cloud Computing (CLOUD), 2011. Sujesha Sudevalayam and Purushottam Kulkarni.
- *Affinity-aware Modeling of CPU Usage with Communicating Virtual Machines.* Journal of Systems and Software (JSS), 2013. Sujesha Sudevalayam, Purushottam Kulkarni.

Thesis Contributions

1. Network-affinity aware CPU Usage Estimation

Prediction of virtualized CPU usage for inter-PM and intra-PM network communication between VMs

2. VM Disk I/O Reduction by Host-cache Manipulation

Reduction of disk I/O by exploiting content similarity within and across virtual machines

VM Disk I/O Reduction by Host-cache Manipulation

- 1 Analysis of existing work (IODEDUP) to show inconsistent performance
- 2 Redirection of I/O requests from within VMs and implicitly manipulate host-cache in content-deduplicated fashion
- 3 Evaluation using public dataset available online
- 4 Case for generation of realistic I/O deduplication benchmarks

Effect of Data Similarity on Host-cache Effectiveness




Figure : Typical virtualized system

Two optimization avenues

- ① Duplicate I/O
- ② Duplicate content in cache

Two orthogonal solutions

- ① I/O deduplication (IODEDUP[1]) :
but **cache inclusiveness problem**
- ② Memory deduplication (Satori[2]) :
dedupes **after** data is fetched



Sources of data similarity

Similar operating systems,
libraries, binaries, file copies

Aim of this work

Improve host-cache effectiveness *using*
I/O deduplication techniques,
i.e., *achieve both* in one stroke.

Traces¹ used for evaluation: Similarity study



Observations

- *homes* & *mail* traces have 95% blocks with sharing factor 1, whereas *webvm* trace has 35% blocks with sharing factor 2
- In *webvm* trace, 45% content occur twice, compared to 6-10% in *homes* and *mail* traces

Conclusions

webvm trace is likely to benefit the most from I/O deduplication

¹Workload traces borrowed from the IODEDUP paper [1]. Traces available online at [3] and SNIA

Existing² I/O deduplication technique: IODEDUP³




Figure : System Architecture of IODEDUP

Functioning

- Creates and maintains content-based cache
- Intercepts read requests & services without accessing disk if possible

²Other related work for I/O deduplication & reduction discussed in thesis.

³*I/O Deduplication: Utilizing Content Similarity to Improve I/O Performance*

Existing² I/O deduplication technique: IODEDUP³




Figure : System Architecture of IODEDUP

Drawbacks

- Content-cache *sizing* needs exploration
- Block-cache still faces *duplicate content* problem

Functioning

- Creates and maintains content-based cache
- Intercepts read requests & services without accessing disk if possible

²Other related work for I/O deduplication & reduction discussed in thesis.

³*I/O Deduplication: Utilizing Content Similarity to Improve I/O Performance*

Existing² I/O deduplication technique: IODEDUP³




Figure : System Architecture of IODEDUP

Drawbacks

- Content-cache *sizing* needs exploration
- Block-cache still faces *duplicate content* problem

Functioning

- Creates and maintains content-based cache
- Intercepts read requests & services without accessing disk if possible

Our contribution

- Perform *study of cache effectiveness* for IODEDUP system, using a custom simulator

²Other related work for I/O deduplication & reduction discussed in thesis.

³*I/O Deduplication: Utilizing Content Similarity to Improve I/O Performance*

Study of cache effectiveness for IODEDUP




Figure : Cache-hit ratios for IODEDUP for *webvm* trace. Total cache 512 MB

Observations

- 1 Read-only trace has lowest performance at content-cache of 10% & highest at 90%
- 2 Write-only performance varies reverse, i.e., highest at 10% and lowest at 90%
- 3 At content-cache setting of 90%, read-only performance is 4× Vanilla, but read/write performance 42% worse than Vanilla.

Conclusion

Inconsistency in achievable cache effectiveness

Fundamental issues preventing efficient I/O reduction

Issues

- 1 IOEDUP system [1] has cache inclusiveness problem
- 2 Memory deduplication [2] works after data is already fetched from disk

Naive solution

- Operate host cache in fully-deduplicated fashion, such that only data not present in cache will ever be fetched from disk

Challenges in implementing naive solution

- 1 Needs change to cache datastructures and/or algorithms
- 2 Needs metadata updates per cache insertion
- 3 Needs invasive monitoring & metadata updates per cache eviction

DRIVE: Using implicit caching hints to achieve disk I/O reduction in virtualized environments

Our approach: Augment the virtual disk driver to **use implicit caching hints** to achieve an approximately fully-deduplicated host cache

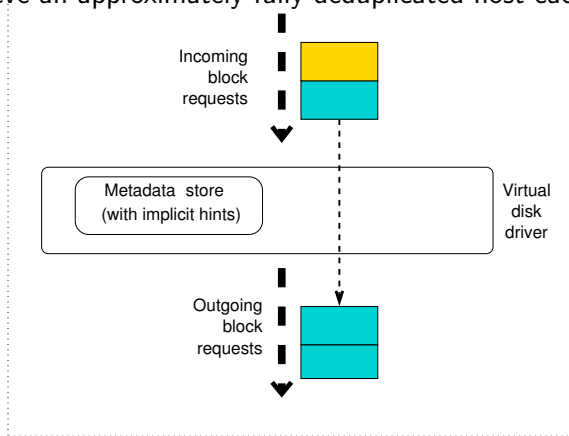


Figure : I/O redirection based on content similarity

Obtaining and using implicit hints for I/O redirection




Figure : Semantics of metadata store.

Obtaining and using implicit hints for I/O redirection




Figure : Semantics of metadata store.

Core idea

- 1 When a block is fetched, it is “known” to be **cached**
- 2 Above is noted in metadata, **marked as leader**
- 3 For next redirection, **leader is used**

Obtaining and using implicit hints for I/O redirection

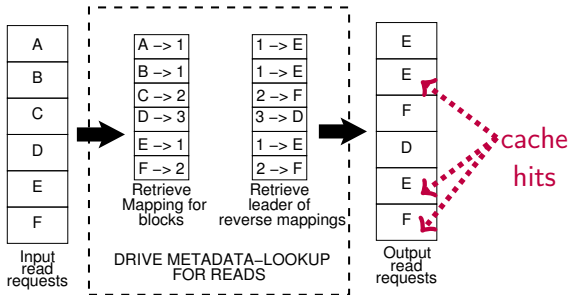



Figure : Semantics of metadata store.

Figure : Example of read request redirection in DRIVE

Core idea

- 1 When a block is fetched, it is "known" to be **cached**
- 2 Above is noted in metadata, **marked as leader**
- 3 For next redirection, **leader is used**

DRIVE: Using implicit caching hints to achieve disk I/O reduction in virtualized environments

System Requirements for DRIVE

- 1 Intercept block read *request* path for **metadata lookup** and I/O redirection, if present
- 2 Intercept block read *return* path for **metadata update**, if not previously present
- 3 Intercept block write *request* path for **metadata invalidation**
- 4 Maintain **implicit caching hints** within metadata to aid efficient I/O redirection.

DRIVE: Using implicit caching hints to achieve disk I/O reduction in virtualized environments

System Requirements for DRIVE


- 1 Intercept block read *request* path for **metadata lookup** and I/O redirection, if present
- 2 Intercept block read *return* path for **metadata update**, if not previously present
- 3 Intercept block write *request* path for **metadata invalidation**
- 4 Maintain **implicit caching hints** within metadata to aid efficient I/O redirection.

Basic difference between DRIVE and IODEDUP

IODEDUP: Services requests using a content-based cache

DRIVE: Redirects I/O requests to duplicate block addresses, to implicitly manipulate host-cache

Block request interception-point for DRIVE



(a) KVM split-driver architecture

(b) Xen split-driver architecture

Interception within VM's front-end driver

- **De-coupling** of the front-end and back-end drivers enables simple I/O redirection
- Results in **implicit** manipulation of host-cache as a content-deduplicated cache
- Exploits individual workload's **content self-similarity**, useful irrespective of co-hosted VMs
- Implementation within generic virtio drivers **obviates dependence** on VMM & guest OS

Evaluating host-cache effectiveness in DRIVE system




Figure : Classification of read responses




Figure : Content deduplication factor of page cache upon *webvm* trace.

Conclusions

- Both *homes* and *mail* workloads have huge number of compulsory misses, whereas the *webvm* workload has significantly fewer.
- DRIVE decreases number of capacity misses to 5% of Vanilla
- DRIVE achieves up to 97% deduplication in block-cache

Identifying similarity in multiple virtual machines




Figure : Read response throughput for aggregated (*homes+webvm*) trace.

Table : Performance for aggregated trace replay

Scheme	Cache-hit ratio (%)	Disk reads reduced(%)	Avg. read response latency (msec)
Vanilla	61.2	1.6	7.9
DRIVE	67.6	18.5	6.5
IODEDUP	62.4	4.3	7.7

Conclusions

- DRIVE completes earlier due to higher number of responses per hour on average⁴.
- Huge margin in percentage of disk reads reduced

Summary of DRIVE

- 1 Performs implicit caching hint-based I/O redirection
- 2 Simulation-based evaluation shows promise—up to 97% content-deduplicated cache achieved
- 3 Further analysis requires more production traces

⁴Throughput derived from measured cache-hits & disk-reads and assumed latency values.

Literature survey for “realistic” dataset generation

Types of datasets generated

- 1 I/O traces (without content) [4, 5, 6, 7, 8]
- 2 Filesystem content (without I/O traces) [9]

Relevant characteristics for I/O traces⁵

Block accessed distribution & Jump distances—*spatial locality*
Run lengths & Block reuse distances—*temporal locality*

General approach

- 1 Capture Multi-dimensional distributions and/or Markov models
- 2 Use above captured models to create new traces with similar properties
- 3 Vary appropriate parameters to create different traces as necessary

⁵*webvm* and *homes* trace characterization presented in thesis.

Content-defined characterization of *webvm* trace




Figure : (a) Block access distribution




Figure : (b) Duplicate block distrib




Figure : (c) Block reuse distribution




Figure : (d) Content reuse distribution

Observations

- Even duplicate content access has **spatial locality** property
- **Temporal locality** is higher for content than block

DRIVE system summary & conclusions

- In this component, we addressed I/O reduction via deduplication
- We analyzed existing work (IODEDUP) and showed that its performance is inconsistent depending on the read/write request-mix of the workload.
- We presented design & implementation of our DRIVE system
- Simulation evaluation shows promise—achieves 97% content deduplication of the host cache.
- We concluded with a survey of publicly available datasets, and benchmark generation literature, to make the case that future work towards I/O deduplication benchmarks is necessary

Publication

DRIVE: Using Implicit Caching Hints to achieve Disk I/O Reduction in Virtualized Environments. Proceedings of the 21st International Conference on High Performance Computing (HiPC), 2014. Sujesha Sudevalayam, Purushottam Kulkarni, Rahul Balani and Akshat Verma.

Bibliography I



Ricardo Koller and Raju Rangaswami.

I/O Deduplication: Utilizing Content Similarity to Improve I/O Performance.

In [Proceedings of the USENIX Conference on File and Storage Technologies \(FAST\)](#), pages 211–224, 2010.



Grzegorz Miłós, Derek G. Murray, Steven Hand, and Michael A. Fetterman.

Satori: Enlightened Page Sharing.

In [Proceedings of the USENIX Annual Technical Conference \(ATC\)](#), pages 1–14, 2009.



Ricardo Koller and Raju Rangaswamy.

Trace: I/O Deduplication: Utilizing Content Similarity to Improve I/O Performance.

Website.

<http://syllab-srv.cs.fiu.edu/doku.php?id=projects:iodedup:start>.



Sriram Sankar and Kushagra Vaid.

Storage Characterization for Unstructured Data in Online Services Applications.

In [Proceedings of the IEEE International Symposium on Workload Characterization \(IISWC\)](#), IISWC '09, pages 148–157. IEEE Computer Society, 2009.



C. Delimitrou, S. Sankar, K. Vaid, and C. Kozyrakis.

Storage I/O Generation and Replay for Datacenter Applications.

In [Proceedings of the IEEE International Symposium on Performance Analysis of Systems and Software \(ISPASS\)](#), pages 123–124, April 2011.



Christina Delimitrou, Sriram Sankar, Kushagra Vaid, and Christos Kozyrakis.

Accurate Modeling and Generation of Storage I/O for Datacenter Workloads, 2011.

Bibliography II



V. Tarasov, S. Kumar, J. Ma, D. Hildebrand, A. Povzner, G. Kuenning, and E. Zadok.

Extracting Flexible, Replayable Models from Large Block Traces.

In [Proceedings of the 10th USENIX Conference on File and Storage Technologies \(FAST\)](#), FAST'12, pages 22–22. USENIX Association, 2012.



Zachary Kurmas, Jeremy Zito, Lucas Trevino, and Ryan Lush.

Generating a Jump Distance Based Synthetic Disk Access Pattern, 2006.



Vasily Tarasov, Amar Mudrankit, Will Buik, Philip Shilane, Geoff Kuenning, and Erez Zadok.

Generating Realistic Datasets for Deduplication Analysis.

In [Proceedings of the USENIX Conference on Annual Technical Conference](#), pages 24–24, 2012.

Thank you! Questions?